

Enhancing Search Precision With Embeddings-Based Semantics

Adarsh Hiremath
NEUROBIO 240: Biological and Artificial Intelligence
ahiremath@college.harvard.edu

May 5, 2023

1. Introduction

Embeddings are numerical representations of words, phrases, or documents that capture their semantic meaning in a vector space. They are commonly used in natural language processing applications, including search, recommendation systems, and sentiment analysis. In search, embeddings can efficiently and accurately retrieve relevant documents or web pages based on the similarity of their semantic content to the user's query.

In December 2022, OpenAI released its newest embedding model, `text-embedding-ada-002`. OpenAI claims that `text-embedding-ada-002` is far more powerful than similar models like BERT. However, little work has been done to quantify the differences between `text-embedding-ada-002` and other machine learning approaches for search. In this paper, I intend to better quantify these differences by doing the following:

- Implementing search with `text-embedding-ada-002`, BERT, and fuzzy matching.
- Creating a custom dataset of my.Harvard classes and their corresponding data.
- Applying those searches to the my.Harvard dataset.
- Quantifying the differences between each search method.
- Building a web application to search Harvard classes and releasing it to the Harvard community.

2. Literature Review

Literature in the field of text embeddings for information retrieval is largely isolated to individual methods, with little work done to compare the effectiveness of different approaches. My paper aims to address this gap by evaluating the performance of three different search methods: `text-embedding-ada-002`, fuzzy matching, and BERT. To inform my evaluation, I draw on the following papers:

- *Text and Code Embeddings by Contrastive Pre-Training* (Neelakantan et al. 2022): an evaluation of unsupervised text embedding models on large-scale semantic search. These researchers measured an improvement of 23.4, 14.7, and 10.6 percent over MSMARCO, Natural Questions and TriviaQA benchmarks.
- *Neural Text Embeddings for Information Retrieval* (Mitra and Craswell 2017): an overview of text embeddings pre-davinci for natural language tasks and information retrieval.
- *Semantic Search With Sentence-BERT for Design Information Retrieval* (Walsh and Andrade 2022): an overview of sentence-BERT, a modification to BERT that uses siamese networks for sentence embeddings.

The Neelakantan paper forms the basis for my `text-embedding-ada002` implementation of semantic search. This paper asserts that it is possible to use a single embedding model to do large-scale semantic search. In my experiments, I use `text-embedding-ada002` as the single embedding model for semantic search and evaluate it on the my.Harvard classes dataset.

The Mitra and Craswell paper provided important context about embeddings in general and was helpful for my general understanding of embeddings.

Finally, the Walsh and Andrade paper was greatly helpful for my initial work on the implementation of the BERT-based search. The insight of passing in the text through a siamese network to determine similarity was particularly helpful.

3. Dataset

Several datasets exist for text classification and aggregation. Initially, I considered testing my semantic search functionality on the following:

- MS MARCO: a text dataset with over 1,000,000 queries and their corresponding search results from the Bing search engine. It is commonly used for passage retrieval.

- *The Reuters Corpus*: a text dataset with over 10,000 articles from Reuters, labeled with categories such as "earnings" and "acquisitions."
- *CORD-19*: a dataset with over 200,000 scholarly articles related to COVID-19 and other coronaviruses. It is commonly used for information retrieval.

However, none of these datasets had queries pre-labeled with "ground truth" about how relevant search responses were. Instead of manually labeling the relevance of search queries for datasets I was unfamiliar with, I opted to use a custom dataset of my.Harvard classes. This dataset contains over 8000 my.Harvard classes organized in the following JSON format:

```
{
  "Id": "GOEnBh8",
  "Subject": "AFRAMER",
  "Number": "11",
  "ClassNumber": "13878",
  "CourseId": "123591",
  "Name": "Introduction to African Studies",
  "Desc": "<p>This course introduces students to the rich diversity and complexity of Africa, including its historical dynamics, economic developments, ...",
  "Profs": "Daniel Agbibo",
  "F22": true,
  "Days": [
    "Th"
  ],
  "M": "N",
  "T": "N",
  "W": "N",
  "R": "Y",
  "F": "N",
  "StartTime": "9:45am",
  "EndTime": "11:45am",
  "HasSection": true,
  "Prereq": "Required of concentrators in African Studies track.",
  "SOC": true,
  "Exam": "12/09/2022 2:00 PM",
  "Location": "Cambridge",
  "URL": "https://locator.tlt.harvard.edu/course/colgsas-123591/2022/fall/13878",
  "AcadOrg": [
    "AAAS",
    "HIST",
    "EMR"
  ]
}
```

Given that my.Harvard's native search only works with ultra-precise search queries for course titles and course names, applying semantic search to my.Harvard classes appeared to be a far more useful and relevant use case for the Harvard community.

my.Harvard has no consistency for the JSON representing each class, making the scraping operation

extremely difficult. For help with generating the dataset, I reached out to the creator of Deez Classes.

4. Search Implementations

Each search technique varies in implementation. `text-embedding-ada-002`, a transformer-based deep learning architecture is utilized along with cosine similarity to measure query and text similarity. The BERT search is implemented using a custom tokenizer, while fuzzy matching uses string-matching algorithms to return results according to a given threshold.

4.1 text-embedding-ada-002

OpenAI's embedding model, `text-embedding-ada-002`, utilizes a transformer-based deep learning architecture to learn embeddings. `text-embedding-ada-002` outperforms the other top models in three standard benchmarks and is useful for working with natural language and code. Semantically similar embeddings are also numerically similar, and the new embeddings endpoint in the OpenAI API provides easy access to the model used to generate embeddings. Although OpenAI offers three families of embedding models - text embeddings text search / similarity, and code search - I only used the base text embedding model for my project.




Figure 1: An example of text embeddings in a vector space.

4.2 BERT

4.3 Fuzzy Matching